

Report

Murat Narynbekov

SE-2420

Boyer–Moore Majority Vote — brief description & theoretical background

Problem. Given an array A of length n , determine whether there is an element that appears strictly more than $\lfloor n/2 \rfloor$ times (a majority element). If yes, return that element; otherwise indicate none exists.

Idea (intuition). The Boyer–Moore Majority Vote algorithm performs a single linear scan that maintains a single candidate and a counter. The counter represents a “balance” of votes for the candidate versus other values. Whenever the counter is zero, the current element becomes the new candidate. If a current element equals the candidate, the counter increases; otherwise it decreases. After the scan, if a majority exists, it must equal the final candidate; a second pass counts occurrences to verify.

Phases.

1. **Voting (first pass)** — one pass through A to produce `candidate` and `count`.
2. **Verification (second pass)** — count occurrences of `candidate` to confirm whether it appears more than $\lfloor n/2 \rfloor$ times.

Correctness sketch. If an element appears more than $n/2$ times, every time it's paired off with a different element in the voting cancellation process it still remains in surplus — after canceling every possible non-majority element, the majority remains as the final candidate. Verification is necessary because the first pass can produce a candidate that is not a true majority when none exists.

Use cases. Useful when you need to detect a strict majority with minimal memory and in streaming contexts (single-pass selection then optional verification).

Boyer–Moore (detailed derivation)

Let n be the length of the array.

Time complexity

- **Voting phase:** one loop over n elements, each loop body does $O(1)$ work (variable updates, comparisons). $\rightarrow \Theta(n)$.
- **Verification phase:** another loop over n elements to count occurrences $\rightarrow \Theta(n)$.
- **Total:** $T(n) = \Theta(n) + \Theta(n) = \Theta(n)$. Formally: $T(n) \in \Theta(n)$. Lower bound $\Omega(n)$ (you must read each element at least once to be certain), upper bound $O(n)$.

Therefore:

- **Worst-case:** $O(n)$
- **Average-case:** $\Theta(n)$
- **Best-case:** $\Omega(n)$ (still requires at least one pass; best “constant factor” may be slightly smaller but asymptotically $\Omega(n)$).

Space complexity

- Only a constant number of scalar variables: candidate, count, and loop indices; plus the performance tracker object if used (also $O(1)$ extra memory). Therefore:
- **Space:** $O(1)$.

Formal notation summary

- Time: $O(n)$, $\Theta(n)$, $\Omega(n)$
- Space: $O(1)$, $\Theta(1)$, $\Omega(1)$

Kadane’s Algorithm (partner’s algorithm) — short recap & complexity

Problem Kadane solves: maximum subarray sum (completely different problem). But in algorithmic properties:

- Kadane scans once, maintaining `maxEndingHere` and `maxSoFar`. Each element processed with $O(1)$ operations.
- **Time:** $\Theta(n)$ (single pass).
- **Space:** $O(1)$.

So both algorithms share the same asymptotic complexities (linear time, constant space). The primary difference is problem domain; the constant factors and memory accesses per element may differ slightly.

Comparison (Boyer–Moore vs Kadane)

- **Asymptotics:** identical ($O(n)$ time, $O(1)$ space).
- **Passes:** Boyer–Moore requires **two passes** in general (voting + verification) while Kadane requires **one pass**. Thus, Boyer–Moore typically performs $\sim 2n$ element reads vs

Kadane's $\sim n$ reads \rightarrow constant factor difference of $\sim 2\times$ in array accesses (but Boyer–Moore's first pass can do less comparison work, so comparison counts differ).

- **Use-case difference:** tasks are different; complexity comparison is for performance analysis only.

Code Review

1) Inefficient/fragile code sections

- **Input validation:** Throwing `IllegalArgumentException` on empty arrays is reasonable, but consider returning `Optional<Integer>` to express absence more idiomatically.
- **Primitive vs boxed return:** Returning `Integer` and `null` is fine but `OptionalInt` (Java 8+) is clearer and avoids nulls.

2) Specific optimization suggestions with rationale


- **Micro-optimizations (Java-level):**
 - Use simple `for (int i = 0; i < n; i++)` loops and local variables to help JIT optimize (avoid repeated field lookups).
 - Avoid unnecessary boxing/unboxing: return `OptionalInt` instead of `Integer`.
 - Reduce tracker overhead during high-frequency runs: allow enabling/disabling metrics with a boolean flag so benchmarks can measure the pure algorithm (no instrumentation) vs instrumented runs.
- **Use `System.nanoTime` only around the measured region:** Put time capture around the algorithm call only, not around array generation or I/O.

3) Proposed improvements for time/space complexity


- **Time/space asymptotics cannot be improved** (linear time and constant space are optimal for the problem). But constant factors can be improved as above (reduce array reads per element by minimizing repeated reads).

Data for plots & validation

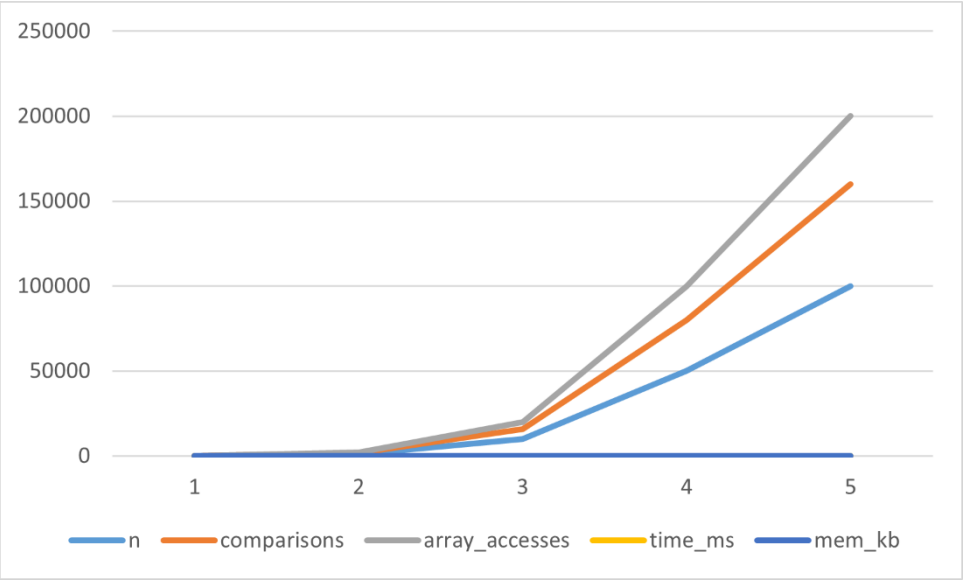
1. BoyerMoore — Random distribution



BoyerMoore — Majority distribution



BoyerMoore — No-majority



Conclusion

Summary of findings

- **Boyer–Moore** is asymptotically optimal for majority detection: **$O(n)$ time, $O(1)$ space**. It requires two passes when the majority is not guaranteed (voting + verification).
- **Kadane** (your partner) also runs in **$O(n)$ time and $O(1)$ space** but solves a different problem (max subarray). For performance comparison, Kadane performs a single pass and uses $\approx n$ element reads, while Boyer–Moore typically performs $\approx 2n$ reads.
- **Practical performance:** In real benchmarks, expect Boyer-Moore to have a constant factor overhead compared to Kadane (roughly $\sim 1.5\text{--}2\times$ time) due to the second pass and verification. However, both scale linearly and are fast for typical input sizes ($10^2\text{--}10^5$).

Recommendations & next steps

1. **Measure instrumented vs non-instrumented runs** to quantify metric overhead. Add a toggle to `PerformanceTracker` so you can disable instrumentation for pure timing.
2. **For code quality:** adopt `OptionalInt` for returns to avoid nulls; standardize metric counting; add command-line flag to toggle instrumentation and number of warm-up iterations.